

# Aspect-based Sentiment Analysis on Hotel Reviews

Nikhilesh Venkat PES1UG23AM192 Rohan Suresh PES1UG23AM240 Prem M Thakur PES1UG23AM214

*"Amazing hotel with amazing staff!"*

Reviewed May 3, 2016

I stayed at The District Boracay this week and had an absolutely wonderful experience. They took care of us from the moment we arrived at the airport, it made our experience extremely pleasurable not having to worry where to go (and we were coming from the farther away airport Kalibo). We had a private boat and air con van to take us directly to the hotel. When we arrived we were welcomed with fresh drinks and cold towels to refresh us. We were then brought to our beautiful room that was large, spacious and modern. We got to experience the spa as well which was the best massage I've gotten in a long time. They have beach chairs out in front of the hotel with a security guard as well (so you can leave all your stuff on your chairs when you want to go swimming and not worry). You have to visit their rooftop restaurant as well, as it is the most stunning view of the sunset and sometimes you have it to yourself as it's not too busy. Positioned in the heart of station 2 you are walking distance to all the shops. Cannot recommend this hotel enough I would come back in a heartbeat!

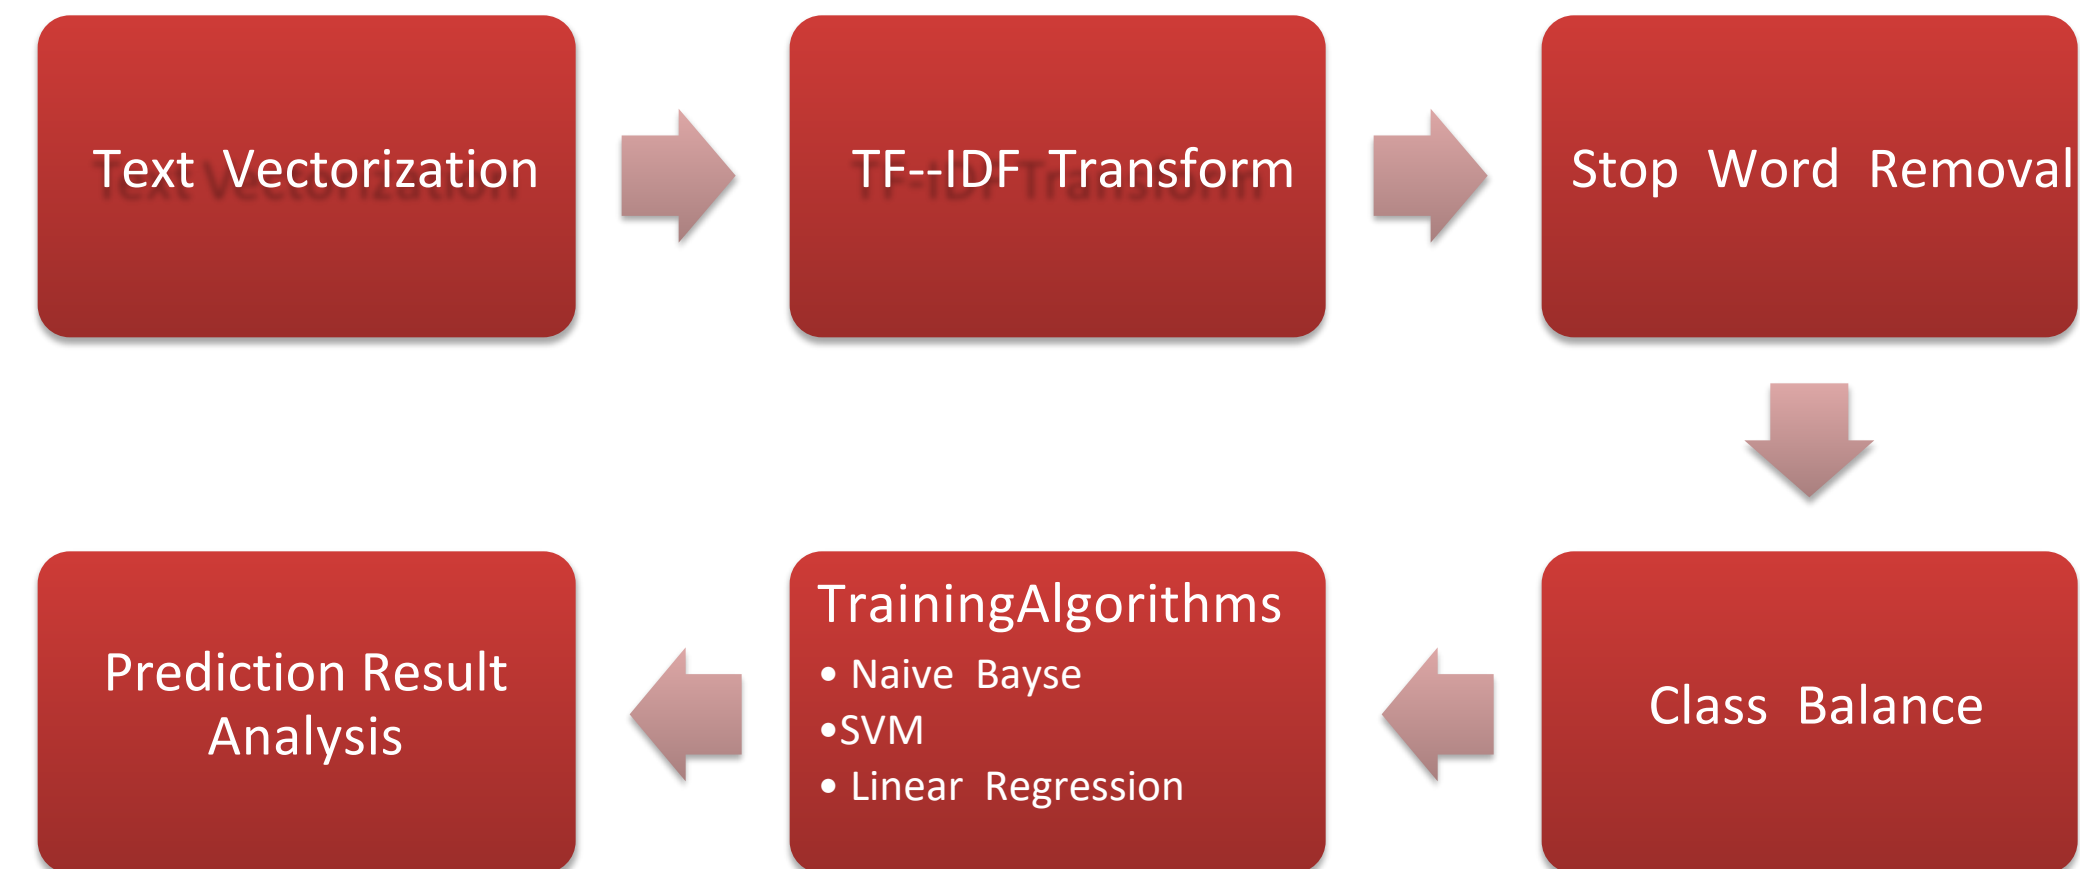
Stayed April 2016, traveled as a couple

Location Rooms Service

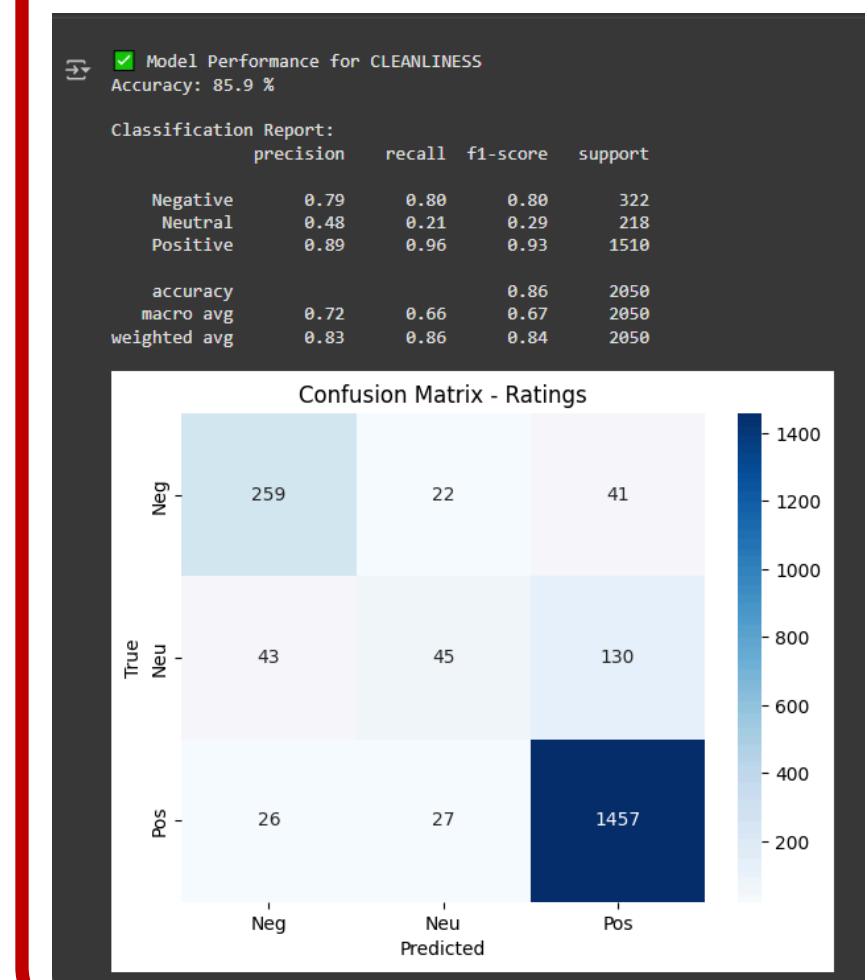
## SAMPLE HOTEL REVIEW AND ASPECT IDENTIFICATION

## Project Objective

Apply supervised learning algorithms on Trip Advisor text reviews. Make prediction on user's opinion on individual aspects.



## FINAL CONFUSION MATRIX



## Handling Data Skewness in Reviews

The TripAdvisor dataset exhibits a class imbalance, with a large majority of reviews being positive. Such skewness can bias the model toward predicting the dominant class, reducing accuracy for neutral and negative reviews.

To address this, we applied SMOTE (Synthetic Minority Over-sampling Technique):

- SMOTE generates synthetic samples for underrepresented classes (Neutral and Negative), balancing the dataset.
- This normalization ensures the model learns features from all classes equally, improving generalization and reducing prediction bias.

Impact:

- Improved performance on minority classes.
- More reliable sentiment predictions across Positive, Neutral, and Negative reviews.

## PROCESS

### Dataset:

- Collected from TripAdvisor hotel reviews.
- Features include text reviews and star ratings.
- Handled class imbalance using SMOTE oversampling.

### Methodology:

1. Data Preprocessing:
  - Removed punctuation and special characters.
  - Converted text to lowercase and removed stopwords.
  - Applied lemmatization to standardize words.
2. Feature Extraction:
  - Used TF-IDF vectorization with n-grams (uni-, bi-, tri-grams) to capture important patterns in text.
3. Model Training:
  - Linear SVM classifier with class balancing.
  - Trained on oversampled data to handle class imbalance.
4. Evaluation:
  - Metrics: Accuracy and F1 Score on train and test sets.
  - Visualized performance with confusion matrices, line plots, and bar charts.

### Key Results:

- High training accuracy indicates the model learned textual patterns effectively.
- Test metrics demonstrate generalization to unseen reviews.
- Sentiment predictions can assist hotels in automated feedback analysis and service improvement.




### Applications:

- Automated review analysis for hotels, restaurants, or e-commerce platforms.
- Customer experience improvement and feedback summarization.
- Real-time sentiment monitoring of online reviews.

### Tools & Libraries:

- Python: Pandas, Numpy, NLTK, Scikit-learn
- Text Processing & ML: TF-IDF, SMOTE, LinearSVC
- Visualization: Matplotlib, Seaborn

## DATASET SKEWNESS TOWARDS POSITIVE REVIEWS



## Model Performance & Improvement Scope

Accuracies & F1 Score:

- The model achieves high training accuracy, indicating effective learning of textual patterns.
- Test accuracy and weighted F1 score demonstrate the model's ability to generalize to unseen reviews.
- Weighted F1 score is particularly useful in imbalanced datasets, as it accounts for class-wise performance, not just overall accuracy.

Improvement Scope:

- Feature Engineering: Incorporate word embeddings (e.g., Word2Vec, BERT) to capture semantic meaning beyond TF-IDF.
- Model Enhancement: Experiment with ensemble methods or deep learning models for better generalization.
- Handling Neutral Reviews: Improve detection of neutral sentiment, which is often ambiguous.
- Cross-Domain Adaptation: Extend to other platforms (e.g., Yelp, Amazon) for broader applicability.

Takeaway:

While the current model performs well, leveraging advanced NLP techniques and richer features can further improve sentiment classification, especially for minority classes.